

The goal of this short note is to explain the relation between two “folklore” results on simple hypothesis testing, and, quite crucially, how they square with each other. Thanks to [Hao-Chung Cheng](#) for illuminating discussions.

For two *fixed* probability distributions $\mathbf{p}, \mathbf{q} \in \Delta(\Omega)$ over a known arbitrary domain Ω , we write $\Psi(\mathbf{p}, \mathbf{q}, \delta)$ for the sample complexity of deciding, with probability of error at most δ , which of these two distributions a sequence of i.i.d. samples from an (unknown) probability distribution $\mathbf{q} \in \{\mathbf{p}_0, \mathbf{p}_1\}$ originates from: specifically, given a uniform prior¹ on $(\mathbf{p}_0, \mathbf{p}_1)$, the error of a test $T: \Omega^n \rightarrow \{0, 1\}$ taking n samples is

$$\delta := \frac{1}{2} \Pr_{\mathbf{p}_0} [T(X_1, \dots, X_n) = 1] + \frac{1}{2} \Pr_{\mathbf{p}_1} [T(X_1, \dots, X_n) = 0] \quad (1)$$

It is well-known (and described in one of these “short notes”) that $\Psi(\mathbf{p}, \mathbf{q}, \delta)$ is characterized by the *squared Hellinger distance* between \mathbf{p} and \mathbf{q} :

Fact 1 (Sample complexity of simple hypothesis testing). *For any $\mathbf{p}_0, \mathbf{p}_1$ and $\delta \in (0, 1]$,*

$$\Psi(\mathbf{p}_0, \mathbf{p}_1, \delta) = \Theta \left(\frac{\log(1/\delta)}{d_H(\mathbf{p}_0, \mathbf{p}_1)^2} \right)$$

where $d_H(\mathbf{p}_0, \mathbf{p}_1) = \frac{1}{\sqrt{2}} \|\sqrt{\mathbf{p}_0} - \sqrt{\mathbf{p}_1}\|_2$ is the Hellinger distance.

Flipping things around, one could ask, given n samples, what the best achievable probability of error δ^* (as in (1)) is as a function of $n, \mathbf{p}_0, \mathbf{p}_1$. Let’s write $\mathcal{E}_n(\mathbf{p}_0, \mathbf{p}_1) := \frac{1}{n} \ln \frac{1}{\delta^*(n, \mathbf{p}_0, \mathbf{p}_1)}$ for this “finite-sample” *error exponent*, so that

$$\delta^* = e^{-n\mathcal{E}_n(\mathbf{p}_0, \mathbf{p}_1)} \quad (2)$$

Then, [Fact 1](#) appears to state that

$$\mathcal{E}_n(\mathbf{p}_0, \mathbf{p}_1) = \Theta \left(d_H(\mathbf{p}_0, \mathbf{p}_1)^2 \right). \quad (3)$$

This is, however, quite annoying, as a classical result in information theory and statistics, the Chernoff bound,² states that $\lim_{n \rightarrow \infty} \mathcal{E}_n(\mathbf{p}_0, \mathbf{p}_1) = C(\mathbf{p}_0, \mathbf{p}_1)$, where

$$C(\mathbf{p}_0, \mathbf{p}_1) := - \min_{\lambda \in [0, 1]} \ln \sum_{x \in \Omega} \mathbf{p}_0(x)^\lambda \mathbf{p}_1(x)^{1-\lambda} \quad (4)$$

is the *Chernoff information* between \mathbf{p}_0 and \mathbf{p}_1 (with the straightforward generalization if Ω is not a discrete domain). *Annoying*, because $C(\mathbf{p}_0, \mathbf{p}_1)$ and $d_H(\mathbf{p}_0, \mathbf{p}_1)^2$ are clearly not the same thing, and having two different (and seemingly *wildly* different) things characterize the same quantity is very confusing at best. So, erm, **how come?**

1 Hellinger distance: (3) is not wrong...

We first provide a self-contained proof of (3); actually, of a stronger version of it, with explicit constants. This is adapting and combining the contents from another of these short notes, “A short note on distinguishing discrete distributions” (2017), and [\[BY02, Theorem 4.7\]](#).

¹One can generalize this to a non-uniform prior; the characterization of the error exponent as Chernoff information will remain the same, as the prior “disappears” asymptotically.

²No, not *that* Chernoff bound.

Lemma 2. For any $\mathbf{p}_0, \mathbf{p}_1$, and $n \geq 1$, we have $e^{2n \ln(1-d_H(\mathbf{p}_0, \mathbf{p}_1)^2)} \leq 2\delta^*(n, \mathbf{p}_0, \mathbf{p}_1) \leq e^{n \ln(1-d_H(\mathbf{p}_0, \mathbf{p}_1)^2)}$, i.e.,

$$-\ln(1-d_H(\mathbf{p}_0, \mathbf{p}_1)^2) + \frac{\ln 2}{n} \leq \mathcal{E}_n(\mathbf{p}_0, \mathbf{p}_1) \leq -2\ln(1-d_H(\mathbf{p}_0, \mathbf{p}_1)^2) + \frac{\ln 2}{n}, \quad (5)$$

which implies (3).

Proof. By the standard interpretation of total variation distance as characterization of the minimal sum of Type I and Type II errors, we have that

$$1 - 2\delta^*(n, \mathbf{p}_0, \mathbf{p}_1) = d_{TV}(\mathbf{p}_0^{\otimes n}, \mathbf{p}_1^{\otimes n})^2 \quad (6)$$

since $\delta^*(n, \mathbf{p}_0, \mathbf{p}_1)$ was defined in (1) as the optimal average error probability when distinguishing \mathbf{p}_0 and \mathbf{p}_1 from n i.i.d. samples. So our task boils down to establishing good enough upper and lower bounds on $d_{TV}(\mathbf{p}_0^{\otimes n}, \mathbf{p}_1^{\otimes n})^2$.

To do so, we will rely on the following two relatively straightforward facts about Hellinger distance, with respect to total variation:

$$1 - \sqrt{1 - d_{TV}(\mathbf{p}_0, \mathbf{p}_1)^2} \leq d_H(\mathbf{p}_0, \mathbf{p}_1)^2 \leq d_{TV}(\mathbf{p}_0, \mathbf{p}_1) \quad (7)$$

and products (tensoring):

$$d_H(\mathbf{p}_0^{\otimes n}, \mathbf{p}_1^{\otimes n})^2 = 1 - \left(1 - d_H(\mathbf{p}_0, \mathbf{p}_1)^2\right)^n. \quad (8)$$

By (8), this implies $d_H(\mathbf{p}_0^{\otimes n}, \mathbf{p}_1^{\otimes n})^2 = 1 - \left(1 - d_H(\mathbf{p}_0, \mathbf{p}_1)^2\right)^n = 1 - e^{n \ln(1-d_H(\mathbf{p}_0, \mathbf{p}_1)^2)}$, and therefore, by (7),

$$d_{TV}(\mathbf{p}_0^{\otimes n}, \mathbf{p}_1^{\otimes n})^2 \geq 1 - e^{n \ln(1-d_H(\mathbf{p}_0, \mathbf{p}_1)^2)} \quad (9)$$

Conversely, from the lower bound from (7) and using (8), we get

$$d_{TV}(\mathbf{p}_0^{\otimes n}, \mathbf{p}_1^{\otimes n})^2 \leq 1 - \left(1 - d_H(\mathbf{p}_0^{\otimes n}, \mathbf{p}_1^{\otimes n})^2\right)^2 = 1 - \left(1 - d_H(\mathbf{p}_0, \mathbf{p}_1)^2\right)^{2n} = 1 - e^{2n \ln(1-d_H(\mathbf{p}_0, \mathbf{p}_1)^2)} \quad (10)$$

and so, combining the two and recalling (6), we finally get

$$e^{2n \ln(1-d_H(\mathbf{p}_0, \mathbf{p}_1)^2)} \leq 2\delta^*(n, \mathbf{p}_0, \mathbf{p}_1) \leq e^{n \ln(1-d_H(\mathbf{p}_0, \mathbf{p}_1)^2)} \quad (11)$$

as claimed. \square

2 ... yet Chernoff is correct.

To square Lemma 2 with the Chernoff bound, which states that

$$\lim_{n \rightarrow \infty} \mathcal{E}_n(\mathbf{p}_0, \mathbf{p}_1) = C(\mathbf{p}_0, \mathbf{p}_1) \quad (12)$$

we need to argue that, while maybe not *equal*, $-\ln(1-d_H(\mathbf{p}_0, \mathbf{p}_1)^2)$ and $C(\mathbf{p}_0, \mathbf{p}_1)$ are always within a factor 2 of each other. Basically, that constant factors *do*, after all, matter.

The first observation is to rewrite $1 - d_H(\mathbf{p}_0, \mathbf{p}_1)^2$ in an equivalent (and standard-ish) form involving the *Bhattacharyya coefficient*,

$$BC(\mathbf{p}_0, \mathbf{p}_1) := \sum_{x \in \Omega} \sqrt{\mathbf{p}_0(x)\mathbf{p}_1(x)}. \quad (13)$$

Namely, one can check that $1 - d_H(\mathbf{p}_0, \mathbf{p}_1)^2 = 1 - BC(\mathbf{p}_0, \mathbf{p}_1)$. This is very convenient, as now we want to compare

$$-\ln(1-d_H(\mathbf{p}_0, \mathbf{p}_1)^2) = -\ln BC(\mathbf{p}_0, \mathbf{p}_1)$$

to

$$C(\mathbf{p}_0, \mathbf{p}_1) = - \min_{\lambda \in [0,1]} \ln \sum_{x \in \Omega} \mathbf{p}_0(x)^\lambda \mathbf{p}_1(x)^{1-\lambda} = - \ln \min_{\lambda \in [0,1]} \sum_{x \in \Omega} \mathbf{p}_0(x)^\lambda \mathbf{p}_1(x)^{1-\lambda}.$$

Getting rid of the logarithms, it would be enough to show that $\min_{\lambda \in [0,1]} \sum_{x \in \Omega} \mathbf{p}_0(x)^\lambda \mathbf{p}_1(x)^{1-\lambda}$ and $\sum_{x \in \Omega} \sqrt{\mathbf{p}_0(x) \mathbf{p}_1(x)}$ are within a quadratic factor of each other. Big if true! And, fortunately, true.

Lemma 3 (Skewed Bhattacharyya coefficients are quadratically related). *For any $\mathbf{p}_0, \mathbf{p}_1$ and $\lambda \in [0, 1]$, we have*

$$\left(\sum_{x \in \Omega} \sqrt{\mathbf{p}_0(x) \mathbf{p}_1(x)} \right)^2 \leq \sum_{x \in \Omega} \mathbf{p}_0(x)^\lambda \mathbf{p}_1(x)^{1-\lambda} \leq 1. \quad (14)$$

In particular, we have

$$\text{BC}(\mathbf{p}_0, \mathbf{p}_1)^2 \leq \min_{\lambda \in [0,1]} \text{BC}_\lambda(\mathbf{p}_0, \mathbf{p}_1) \leq \text{BC}(\mathbf{p}_0, \mathbf{p}_1) \quad (15)$$

where $\text{BC}_\lambda(\mathbf{p}_0, \mathbf{p}_1) = \sum_{x \in \Omega} \mathbf{p}_0(x)^\lambda \mathbf{p}_1(x)^{1-\lambda}$ denotes the λ -skewed Bhattacharyya coefficient.

Proof. Fix any $\lambda \in (0, 1)$ (the cases $\lambda \in \{0, 1\}$ being immediate). First, we have

$$\sum_{x \in \Omega} \mathbf{p}_0(x)^\lambda \mathbf{p}_1(x)^{1-\lambda} = \sum_{x \in \Omega} \mathbf{p}_1(x) \cdot \left(\frac{\mathbf{p}_0(x)}{\mathbf{p}_1(x)} \right)^\lambda \leq \left(\sum_{x \in \Omega} \mathbf{p}_1(x) \cdot \frac{\mathbf{p}_0(x)}{\mathbf{p}_1(x)} \right)^\lambda = \left(\sum_{x \in \Omega} \mathbf{p}_0(x) \right)^\lambda = 1$$

using Jensen's inequality for the concave function $x \mapsto x^\lambda$.

Second, let's use Hölder's (the generalized version, with 3 vectors) with exponents $(2/(1-\lambda), 2, 2/\lambda, 2)$, which satisfy $\frac{1-\lambda}{2} + \frac{1}{2} + \frac{\lambda}{2} = 1$. We have

$$\begin{aligned} \sum_{x \in \Omega} \mathbf{p}_0(x)^{1/2} \mathbf{p}_1(x)^{1/2} &= \sum_{x \in \Omega} \mathbf{p}_0(x)^{\frac{1-\lambda}{2}} \cdot \mathbf{p}_0(x)^{\frac{\lambda}{2}} \mathbf{p}_1(x)^{\frac{1-\lambda}{2}} \cdot \mathbf{p}_1(x)^{\frac{\lambda}{2}} \\ &\leq \left(\sum_{x \in \Omega} \mathbf{p}_0(x) \right)^{\frac{1-\lambda}{2}} \left(\sum_{x \in \Omega} \mathbf{p}_0(x)^\lambda \mathbf{p}_1(x)^{1-\lambda} \right)^{\frac{1}{2}} \left(\sum_{x \in \Omega} \mathbf{p}_1(x) \right)^{\frac{\lambda}{2}} \quad (\text{Hölder}) \\ &= \left(\sum_{x \in \Omega} \mathbf{p}_0(x)^\lambda \mathbf{p}_1(x)^{1-\lambda} \right)^{\frac{1}{2}}, \end{aligned}$$

concluding the proof. □

This readily implies that, for every $\mathbf{p}_0, \mathbf{p}_1$,

$$-\ln(1 - d_H(\mathbf{p}_0, \mathbf{p}_1)^2) \leq C(\mathbf{p}_0, \mathbf{p}_1) \leq -2\ln(1 - d_H(\mathbf{p}_0, \mathbf{p}_1)^2) \quad (16)$$

and sanity is restored.

References

- [BY02] Ziv Bar-Yossef. *The Complexity of Massive Data Set Computations*. PhD thesis, UC Berkeley, 2002. Adviser: Christos Papadimitriou. Available at http://webee.technion.ac.il/people/zivby/index_files/Page1489.html.